




The graph constrains the predictive distribution during training; every prediction retains the graph-induced prior that informed it.